

32-bit Modular Arithmetic and SIMD Vectorization

Authors: Duc Vinh Nguyen, Henry Zheng

Supervisor: Vincent Neiger, Hugo Passe

Git: <https://github.com/hertz24/pcca7.git>

May 2026



Contents

1	Introduction	2
2	Fundamental Cost of Modular Reduction	2
3	Shoup's Modular Multiplication	2
3.1	Input Constraints	2
3.2	Theorem	3
3.3	Input Hypotheses	3
3.4	Precomputation Logic (Barrett Reduction)	4
3.5	Fundamental Principle	4
3.6	Proof	4
3.7	Advantages	6
3.8	SIMD Vectorization Implementation	6
3.8.1	AVX2 and AVX-512	6
3.8.2	NEON	7
4	Performance	8
4.1	AMD Overall Hypothesis	8
4.2	NEON Overall Hypothesis	8
4.3	Equipment Details for AMD Architecture	8
4.4	Equipment Details for ARM Architecture	8
4.5	AMD: Choice of Unrolling	8
4.6	AMD: Choice Between <code>mullo_epi32</code> and <code>mul_epu32</code>	10
4.7	AMD: Overall Performance	12
4.8	NEON's Choice of Unrolling	13
4.9	NEON's Overall Performance	15
5	Overall Conclusion	16
6	Improvements	17
A	SIMD Implementations	18
A.1	AVX2	18
A.2	AVX-512	21
A.3	NEON	25

1 Introduction

This project explores and optimizes modular reductions modulo primes $p < 2^{32}$ using SIMD vectorization (AVX2, AVX-512, NEON) along with various techniques and algorithms.

Objective: The aim of this study is to develop an optimized implementation of modular reduction for 32-bit primes, minimizing execution time through SIMD vectorization.

2 Fundamental Cost of Modular Reduction

The naive approach to computing $a \pmod n$ is to evaluate:

$$a - \left\lfloor \frac{a}{n} \right\rfloor \cdot n \quad (1)$$

However, on modern processor architectures (Intel, AMD, ...), integer division has a significantly higher computational cost compared to basic arithmetic operations such as additions and multiplications.

Since equation (1) requires an integer division, a multiplication and a subtraction operation, it is considered highly inefficient from both a software and hardware perspective.

3 Shoup's Modular Multiplication

Shoup's algorithm is based on precomputation to avoid costly operations. His work was first appeared directly in the source code of NTL (Number Theory Library)[5]

Algorithm 1: Shoup's modular multiplication algorithm

Input : $A, B \in [0, p[$, we suppose that $p < \beta/2$
precompute $B' = \lfloor B\beta/p \rfloor$

Output: $R = AB \pmod p$

```
1  $Q \leftarrow \lfloor AB'/\beta \rfloor$  ;  
2  $R \leftarrow (AB - Qp) \pmod \beta$  ;  
3 if  $R \geq p$  then  
4    $R \leftarrow R - p$ ;
```

3.1 Input Constraints

To ensure the correctness of the modular reduction and prevent register overflow during SIMD operations, the following conditions must be satisfied:

- **Word Size (β):** Set to 2^{32} for 32-bit modular arithmetic
- **Input 1 (A):** An integer defined over 32 bits
- **Input 2 (B):** Integer within the range $[0, p[$
- **Modulus (p):** A prime number defined over 31 bits ($p < 2^{31}$)

3.2 Theorem

Let p be a prime number such that $p < \beta/2$ where $\beta = 2^{32}$. Given an integer A defined over 32 bits and $B \in [0, p)$, let B' be the precomputed value $B' = \lfloor B\beta/p \rfloor$. If the estimated quotient is defined as $Q = \lfloor AB'/\beta \rfloor$, then the intermediate remainder $R = AB - Qp$ satisfies the following condition:

$$0 \leq AB - Qp < 2p$$

3.3 Input Hypotheses

Can p be any integer over 32 bits?

No, one particularly interesting case where Shoup's algorithm will fail is for $p \geq 2^{31}$. Since we know that:

$$0 \leq AB - Qp < 2p$$

If $p \geq 2^{31}$, then $2p \geq 2^{32}$. Therefore, in some cases $AB - Qp$ may exceed 2^{32} . This causes tremendous errors in the calculation because all bits beyond 32 bits are discarded when we reduce modulo β .

Does p necessarily have to be a prime number?

No, the algorithm mainly uses precomputations (Barrett Reduction), optimized for calculation of $a \bmod n$ without constraints on n . But for the sake of this project, n is a prime number (p).

Can b be any integer over 32 bits?

No, suppose $b > p$, we can rewrite b as:

$$b = q' \cdot p + r'$$

Thus the calculation for the precomputed value B' is:

$$B' = \left\lfloor \frac{(q' \cdot p + r') \cdot 2^{32}}{p} \right\rfloor$$

$$B' = q' \cdot 2^{32} + \left\lfloor \frac{r' \cdot 2^{32}}{p} \right\rfloor$$

With a focus on 32 bits in this project, B' in the case of $b > p$ cannot be stored within 32 bits and will overflow.

Can a be any integer over 32 bits?

Yes, a can be any integer defined over 32 bits.

3.4 Precomputation Logic (Barrett Reduction)

The precomputed value B' is defined over 32 bits as:

$$B' = \lfloor (B \cdot 2^{32})/p \rfloor$$

When computing $A \cdot B \pmod{p}$:

$$A \cdot B \pmod{p} = A \cdot B - \left\lfloor \frac{A \cdot B}{p} \right\rfloor \cdot p$$

We can rewrite this:

$$\begin{aligned} A \cdot B \pmod{p} &= A \cdot B - \left\lfloor \frac{A \cdot \left\lfloor \frac{B \cdot 2^{32}}{p} \right\rfloor}{2^{32}} \right\rfloor \cdot p \\ &\approx A \cdot B - \left\lfloor \frac{A \cdot B'}{2^{32}} \right\rfloor \cdot p \end{aligned}$$

3.5 Fundamental Principle

Given x , the quotient of two 32-bit integers:

The proof of this algorithm is based on the non-decimal value of $\lfloor x \rfloor$:

$$x - 1 < \lfloor x \rfloor \leq x$$

We can rewrite this as:

$$0 \leq x - \lfloor x \rfloor < 1$$

3.6 Proof

This following proof can be found in "Faster arithmetic for number-theoretic transforms" page 3 [3]

The algorithm uses the precomputation $B' = \lfloor B\beta/p \rfloor$ such that:

$$\begin{aligned} 0 &\leq \frac{B\beta}{p} - \left\lfloor \frac{B\beta}{p} \right\rfloor < 1 \\ 0 &\leq \frac{B\beta}{p} - B' < 1 \end{aligned} \tag{2}$$

Using the same logic, the algorithm estimates the quotient $Q = \lfloor AB'/\beta \rfloor$ such that:

$$\begin{aligned} 0 &\leq \frac{AB'}{\beta} - \left\lfloor \frac{AB'}{\beta} \right\rfloor < 1 \\ 0 &\leq \frac{AB'}{\beta} - Q < 1 \end{aligned} \tag{3}$$

We now prove that the remainder $R = (A \cdot B - Q \cdot p)$ lies within the interval $[0, 2p[$:

We first multiply the equation (1) by Ap/β :

$$\begin{aligned} 0 &\leq \left(\frac{B\beta}{p} - B' \right) \cdot \frac{Ap}{\beta} < 1 \cdot \frac{Ap}{\beta} \\ 0 &\leq AB - \frac{AB'p}{\beta} < \frac{Ap}{\beta} \end{aligned} \tag{4}$$

Similarly, we multiply the equation (2) by p :

$$\begin{aligned} 0 &\leq \left(\frac{AB'}{\beta} - Q \right) \cdot p < 1 \cdot p \\ 0 &\leq \frac{AB'p}{\beta} - Qp < p \end{aligned} \tag{5}$$

By adding equations (3) and (4), we deduce:

$$0 \leq (AB - \cancel{\frac{AB'p}{\beta}}) + (\cancel{\frac{AB'p}{\beta}} - Qp) < \frac{Ap}{\beta} + p$$

Simplifying this result, we are left with:

$$\begin{aligned} 0 &\leq AB - Qp < p + \frac{Ap}{\beta} \\ 0 &\leq AB - Qp < p \cdot \left(1 + \frac{A}{\beta} \right) \end{aligned}$$

By definition, $A \in [0, p[$ and $p < \beta$, thus $A < \beta$. Therefore we can write:

$$\begin{aligned} 0 &\leq AB - Qp < p \cdot \left(1 + \frac{A}{\beta} \right) < p \cdot \left(1 + \frac{\beta}{\beta} \right) \\ 0 &\leq AB - Qp < p \cdot (1 + 1) \\ 0 &\leq AB - Qp < 2p \\ 0 &\leq R < 2p \end{aligned}$$

Conclusion: We have now proven that $R \in [0, 2p[$. The final if statement of the algorithm helps correct the remainder: if $R \geq p$, one more step $R \leftarrow R - p$ is needed; if not, R is the correct result.

3.7 Advantages

Using the precomputed value B' to approximately calculate $\left\lfloor \frac{A \cdot B}{p} \right\rfloor$ improves upon the naive method in two ways:

1. **Storing the precomputed value:** B' is computed with one division and can be stored for use in other modular reductions, provided the same B and the same p are used.
2. **Efficient operations:** Shoup's algorithm replaces slow division operations with fast multiplications and bit shifts. The expression $\frac{A \cdot B'}{2^{32}}$ requires only one multiplication operation and a right bit shift by 32 bits. On a CPU, both multiplication (≈ 0.33 – 0.67 cycles per operation) and bit shift (≈ 0.5 – 1 cycles per operation) are significantly more efficient than division (≈ 6 cycles per operation). A division is approximately 4–8 times more costly than a multiplication and bit shift together. [1]

3.8 SIMD Vectorization Implementation

We decided to create three versions for implementing a high-performance Shoup algorithm: AVX2, AVX-512 and NEON.

3.8.1 AVX2 and AVX-512

For the AVX version, we decided to use two main built-in AVX2 functions that can be found on Intel's Intrinsic Guide: [4]

- `_mm256_mul_epu32(...)` : This function treats the input as unsigned 32-bit integers. However, it only multiplies the even-indexed slots (0, 2, 4, 6) of the vector to make room for the full 64-bit results.
- `_mm256_mullo_epi32(...)` : This function treats the input as signed 32-bit integers. It multiplies all 8 pairs of numbers but only stores the bottom half of the result.

Variant 1: Intel `_mm256_mul_epu32(...)` version: (Listings 1):

This implementation uses `_mm256_mul_epu32(...)` at all stages of calculation, so we must shuffle the data registers into their correct positions. The following are the steps per loop:

- Load 8 32-bit integers into the register
- Split the loaded register into even-indexed slots (0, 2, 4, 6) and odd-indexed slots (1, 3, 5, 7) into 64-bit registers
- Compute Shoup's modular reduction (`_shoup_avx2(...)`) over these registers
- Recombine the even and odd indexed slots and store
- Compute the leftover elements

Variant 2: Intel `_mm256_mullo_epi32(...)` version 1: (Listings 2):

This implementation uses `_mm256_mul_epu32(...)` to compute the quotient:

$$Q = \lfloor (A \cdot B') / 2^{32} \rfloor$$

but `_mm256_mullo_epi32()` to compute all other steps per loop. The following are the steps per loop:

- Load 8 32-bit integers into the register
- Split the loaded register into even-indexed slots (0, 2, 4, 6) and odd-indexed slots (1, 3, 5, 7) into 64-bit registers
- Compute Shoup’s modular reduction (`_shoup_mullo_avx2(...)`) over these registers
- Recombine the even and odd indexed slots and store
- Compute the leftover elements

Variant 3: Intel `_mm256_mullo_epi32(...)` version 2: (Listings 3):

Derived from the previous version, this version aims to regroup the vector registers immediately after the multiplication step:

$$Q = \lfloor (A \cdot B') / 2^{32} \rfloor$$

By restoring the standard 8-lane 32-bit layout early, the function eliminates the need to perform parallel logic paths. This, in theory, should reduce the workload.

In order to implement this version, we create an auxiliary function `_mulhi_emul_avx2(...)` (see Listing 3) which emulates a native “multiply-high” instruction for unsigned 32-bit integers. It computes the full 64-bit products of two packed vectors, extracts the most significant 32 bits from each result, and repacks them into a single 256-bit destination register.

The following are the steps per loop:

- Load 8 32-bit integers into the register
- Split the loaded register into even-indexed slots (0, 2, 4, 6) and odd-indexed slots (1, 3, 5, 7) into 64-bit registers to compute the quotient
- Recombine the even and odd indexed slots
- Continue with Shoup’s modular reduction (`_shoup_mullo_v2_avx2(...)`)
- Compute the leftover elements

3.8.2 NEON

Variant NEON: TODO

4 Performance

The performance benchmarks were run on AMD and ARM architectures, by executing `./pcca7 -b.bits 31 -pts 1000`.

All execution time were measured using `prof_repeat()` from the FLINT library [2]

Note: For NEON, `prof_repeat()` took too much time (8 hours for 1 test) thus we had to rewrite this function natively

4.1 AMD Overall Hypothesis

AVX With a rough preliminary analysis, using SIMD operations should provide us with a maximum acceleration of 8x–16x since the registers can hold and operate on 8 32-bit unsigned integers (for AVX2) and 16 32-bit unsigned integers (for AVX-512) simultaneously.

One caveat that must be considered is that all implementations use `_mm256_mul_epu32(...)` to calculate the quotient defined as $Q = \lfloor A \cdot B' / \beta \rfloor$. This operation requires a data shuffle to correctly align the vectors into 64-bit lanes and then recombine them at later stages. This reduces the acceleration factor of SIMD implementations, at worst, by half (4x–8x depending on AVX2 or AVX-512).

4.2 NEON Overall Hypothesis

NEON uses 128-bit wide registers, This means a single vector can hold up to four 32-bit unsigned integers simultaneously. In consequence, using SIMD operation, with rough preliminary analysis should provide us with a maximum acceleration of 4x.

Just like our AVX implementation, all implementation in NEON require the use of `vmul_u32(...)` (`_mm256_mul_epu32(...)` equivalence in NEON). Thus by the same reasoning, this cuts the acceleration factor of the SIMD implementation by at worse, by half (2x).

4.3 Equipment Details for AMD Architecture

- Processor: AMD Ryzen™ 5 8600G (Zen 4)
- Instruction Sets Supported: AVX, AVX2 and AVX-512
- Compiler: GCC 13.3.0 with optimization flag `-O3` and `-march=native`

Note: AMD Zen 4 simulates AVX-512 using 2 register of AVX2. This is not a real AVX-512

4.4 Equipment Details for ARM Architecture

- Processor: Apple M1
- Instruction Sets Supported: NEON
- Compiler: GCC 13.3.0 with optimization flag `-O3`

Note: The benchmarks were run in a virtual machine.

4.5 AMD: Choice of Unrolling

One way to optimize for faster functions through vectorization is loop unrolling. This technique consists of expanding the body of a loop to process multiple iterations at once.

The results below were obtained by measuring the execution time between a standard version and an unrolled version of Variant 1 (Listing 1) with AVX2 optimisation.

Hypothesis: The unrolled version should perform better than the standard version.




Figure 1: Execution time (in microseconds) comparison: Standard vs. Unrolled (AVX)

Size	Standard version	Unrolled version	Percentage difference
100	0.034	0.037	+8.82%
200	0.056	0.056	+0.00%
300	0.076	0.076	+0.00%
400	0.098	0.090	-8.12%
500	0.116	0.107	-7.76%
600	0.131	0.124	-5.34%
700	0.154	0.142	-7.80%
800	0.190	0.155	-18.42%
900	0.188	0.176	-6.38%
1000	0.205	0.191	-6.83%

Table 1: Execution time of implementations with and without unrolling in microseconds (AVX)

Observation: From the graph and the results table, unrolling have little to no benefit for smaller vector sizes of (under 400). For computing bigger vector sizes, we gradually

observe a slight performance increase for the Unrolled version approximately 6-8%

Conclusion: The results obtained align with our preliminary assumption that applying unrolling to the loops will increase performance. Although the speedup is only around 5% (x1.05 speedup factor), we can still conclude that unrolling is favorable for our implementation of Shoup’s modular multiplication

4.6 AMD: Choice Between `mullo_epi32` and `mul_epu32`

`mullo_epi32` and `mul_epu32` both produce at most 64-bit results. The key difference is that the former truncates the result and stores only the lower 32 bits, whereas the latter stores the full 64-bit result.

The results below were obtained by measuring the execution time of the three variants with unrolling described in Section 3.8.1 using AVX2 optimisation.

Hypothesis: The two `mullo_epi32` variants (Variant 2 and 3) should perform better than the `mul_epu32` variant (Variant 1) because they use fewer operations overall. The Variant 3 should outperform Variant 2 .




Figure 2: Execution time (in microseconds) comparison of all implementation variants (AVX)

Size	Variant 1	Variant 2	Acceleration	Variant 3	Acceleration
100	0.037	0.037	+0.00%	0.024	-35.13%
200	0.056	0.056	+0.00%	0.045	-19.64%
300	0.076	0.077	+1.32%	0.057	-25.00%
400	0.090	0.087	-3.33%	0.070	-22.22%
500	0.107	0.116	+8.41%	0.081	-24.29%
600	0.124	0.127	+2.41%	0.089	-28.22%
700	0.142	0.147	+3.52%	0.102	-28.17%
800	0.155	0.161	+3.87%	0.113	-27.10%
900	0.176	0.182	+3.40%	0.123	-30.11%
1000	0.191	0.196	+2.61%	0.130	-31.93%

Table 2: Execution time of all implementations and the speedup in comparison to `mul_epu32` in microseconds (AVX)

Observation: From the graph, we can see that for small vector sizes (under 500), the performance difference between Variant 1 and Variant 2 is negligible. However, for larger vector sizes, Variant 1 actually outperforms slightly.

As for Variant 3, it outperforms both previous variants for all sizes.

On average, Variant 1 is 2.22% slower than Variant 2, whereas Variant 3 is 24.85% faster than Variant 1.

Conclusion: The measured results contradict our hypothesis that both `mullo_epi32` variants are faster than `mul_epu32`. Even though Variant 2 should in theory outperform the Variant 1 because 8 integers are processed per iteration. In practice, Variant 2 performs worse because its SIMD operation cost is more cycles (1 cycle per operation `mullo_epi32`) compared to Variant 1 (0.5 cycle per operation `mul_epu32`) according to `uops.info` [1]

Although Variant 2 is a worse implementation of Shoup’s modular multiplication, we can still conclude that Variant 3 is the best implementation among the variants, with a performance gain of approximately 25% (x1.25 speedup factor).

4.7 AMD: Overall Performance

It is worth noting that AVX2 uses 256-bit registers to process up to eight 32-bit integers, while AVX-512 uses 512-bit registers for up to sixteen. Furthermore, the AVX-512 version uses the same functions and operations as the AVX2 version.

To measure the best performance output, we decided that the fastest implementation : Variant 3 (see Listing 3) with unrolling

The results below were obtained by measuring the execution time of different implementations of modular reduction:

- **Naive scale:** The base operation (%)
- **Shoup scale (reference) :** Shoup’s modular multiplication without optimization
- **Shoup scale (FLINT) :** FLINT’s modular reduction implementation
- **Shoup scale (AVX2) :** Shoup’s modular multiplication with AVX2 optimization (Variant 3 with unrolling)
- **Shoup scale (AVX-512) :** Shoup’s modular multiplication with AVX-512 optimization (Variant 3 with unrolling)




Figure 3: Execution time comparison between different implementations of modular reduction (AVX)

Size	Naive	Shoup ref	FLINT	Shoup AVX2	Shoup AVX-512
100	0.187	0.135	0.098	0.024	0.022
200	0.397	0.277	0.193	0.045	0.035
300	0.594	0.408	0.279	0.056	0.048
400	0.794	0.546	0.365	0.070	0.053
500	0.986	0.687	0.454	0.081	0.068
600	1.186	0.806	0.541	0.089	0.075
700	1.382	0.940	0.621	0.102	0.091
800	1.574	1.068	0.709	0.113	0.094
900	1.786	1.198	0.794	0.123	0.101
1000	1.971	1.330	0.886	0.130	0.116

Table 3: Execution time for different implementations in relation to the size of the vector in microseconds (AVX)

Algorithm	Average %	Acceleration Factor
FLINT	100.00%	1x
Shoup (AVX2)	16.86%	5.9x
Shoup (AVX-512)	14.35%	7.0x

Table 4: Average performance summary of FLINT in relation to Shoup AVX2 and Shoup AVX-512 implementations (AVX)

Observation: From the results retrieved, we want to primarily study the speedup factor of our Shoup’s modular multiplication using AVX2 and AVX-512 compared to FLINT’s implementation. Shoup (AVX2) is roughly 5.9x faster than FLINT and Shoup (AVX-512) is 7.0x faster.

Conclusion: Our Shoup’s modular multiplication with AVX2 and AVX-512 optimizations adapted to 32 bits is 5–7 times faster than FLINT. These results are consistent with our overall hypothesis.

4.8 NEON’s Choice of Unrolling

The results below were obtained by measuring the execution time between a standard version and an unrolled version of Variant 1 (NEON) (Listing 7).

Hypothesis: The unrolled version should perform better than the standard version.




Figure 4: Execution time (in microseconds) comparison: Standard vs. Unrolled (NEON)

Size	Standard version	Unrolled version	Percentage difference
100	0.000289	0.000307	+6.23%
200	0.000637	0.000622	-3.35%
300	0.000858	0.000876	+2.10%
400	0.001076	0.001096	+1.86%
500	0.001307	0.001354	+3.60%
600	0.001491	0.001543	+3.49%
700	0.001737	0.001783	+2.65%
800	0.001926	0.002035	+5.66%
900	0.002158	0.002283	+5.79%
1000	0.002369	0.002519	+6.33%

Table 5: Execution time of implementations with and without unrolling in microseconds (NEON)

Observation: From the results table, loop unrolling offers little to no benefit for smaller vector sizes (under 400), with the exception of a minor anomaly at size 200. As the vector size increases beyond this point, we gradually observe a slight performance degradation for the unrolled version, which executes approximately 3% to 6% slower than the standard version with bigger vector sizes.

Conclusion: The results obtained contradict our preliminary assumption that applying

unrolling to the loops would increase performance. Instead of a speedup, the unrolled implementation introduced a 1.03x to 1.06x slowdown factor for larger inputs. Consequently, we can conclude that loop unrolling is unfavorable for our specific implementation of Shoup’s Modular Multiplication on NEON.

4.9 NEON’s Overall Performance

The results below were obtained by measuring the execution time of different implementations of modular reduction:

- **Naive scale:** The base operation (%)
- **Shoup scale (reference) :** Shoup’s modular multiplication without optimization
- **Shoup scale (FLINT) :** FLINT’s modular reduction implementation
- **Variant 1 (NEON) :** See Listing 7 (without unrolling)
- **Variant 2 (NEON) :** See Listing 8 (without unrolling)




Figure 5: Execution time comparison between different implementations of modular reduction (NEON)

Size	Naive	Shoup ref	FLINT	Variant 1 (NEON)	Variant 2 (NEON)
100	0.000514	0.000692	0.000653	0.000289	0.000279
200	0.001099	0.001274	0.001310	0.000637	0.000559
300	0.001573	0.001848	0.001871	0.000859	0.000763
400	0.002056	0.002389	0.002380	0.001077	0.000954
500	0.002542	0.003021	0.002929	0.001308	0.001162
600	0.003034	0.003422	0.003477	0.001492	0.001360
700	0.003630	0.003939	0.004094	0.001737	0.001733
800	0.003993	0.004420	0.005752	0.001927	0.001777
900	0.004488	0.005048	0.005128	0.002159	0.002000
1000	0.005093	0.005585	0.005683	0.002370	0.002208

Table 6: Execution time for different implementations in relation to the size of the vector in microseconds (NEON)

Algorithm	Average %	Acceleration Factor
FLINT	100.00%	1x
Variant 1	41.63%	2.4x
Variant 2	38.45%	2.6x

Table 7: Average performance summary of FLINT in relation to the Variant 1 (NEON) and Variant 2 (NEON)

Observation: From the results table, we firstly notice irregularities between the naive methode and FLINT. The library FLINT on NEON is actually slower.

Even so, both variants implemented with NEON are overall much faster than their FLINT counter part with Variant 1 being 2.4x faster than FLINT’s version. Variant 2 being 2.6 times faster.

Conclusion: he acceleration factors observed for both variants are consistent with the NEON theoretical hypothesis outlined in Section 4.2.

However, despite these relevant speedups, the results should be interpreted with caution. There is a noticeable lack of consistency in execution time when comparing FLINT’s optimized implementation against the naive modulo (%) operation.

5 Overall Conclusion

For AVX architectures, we researched and implemented several versions of Shoup’s modular reduction that outperforms FLINT’s implementation. **Variant 3: Intel `_mm256_mullo_epi32(...)`** **version 2** is the best performing version, with an acceleration factor of 5.9x for AVX2 and 7.0x for AVX-512.

For NEON architecture, **Variant 2: `_shoup_mullo_neon`** is the best performing implmentation with an acceleration of 2.6x

Note: Results from NEON could be misleading as all of the results are of order 10^{-4} microseconds meanwhil results from AVX are of order 10^{-1} microseconds. Suppose these

results are corrects, the naive operation (%) would be 1000 times faster than any implementation in AVX2 and AVX-512

6 Improvements

Testings on INTEL: : These results were obtained on AVX architecture built on AMD (Zen 4). Further test and analysis needs to be carried out on the AVX architecture build on Intel as well.

Extensive testing on NEON: : The results obtained for NEON are inconsistance and does not reflect the theoretical analysis. More test should be carried out on different NEON machine architecture to ensure a better scope of speedup.

References

- [1] Andreas Abel and Jan Reineke. *uops.info: Instruction Latency and Throughput Tables*. 2018.
- [2] William Hart et al. *FLINT: Fast Library for Number Theory*. Version 3.5.0, <https://flintlib.org>.
- [3] David Harvey. “Faster arithmetic for number-theoretic transforms”. In: *CoRR* abs/1205.2926 (2012). arXiv: [1205.2926](http://arxiv.org/abs/1205.2926). URL: <http://arxiv.org/abs/1205.2926>.
- [4] Intel. *Intel® Intrinsics Guide*. Version 3.6.9. 2024. URL: <https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html>.
- [5] Victor Shoup. *NTL: A Library for Doing Number Theory*. <https://shoup.net/ntl/>. version 11.6.0; Release date: 2025.11.07. 2005.

A SIMD Implementations

A.1 AVX2

Listing 1: Variant 1 (AVX2)

```
1 static inline __m256i _shoup_avx2(__m256i va, __m256i vb, __m256i
  vb_precomp, __m256i vp)
2 {
3     // vq = [ (a_0 * b_precomp_0 = q_0q_1), (a_2 * b_precomp_2 = q_2q_3),
4         (a_4 * b_precomp_4 = q_4q_5), (a_6 * b_precomp_6 = q_6q_7) ]
5     __m256i vq = _mm256_mul_epu32(va, vb_precomp);
6
7     // vq = [ (0, q_0), (0, q_2), (0, q_4), (0, q_6) ]
8     vq = _mm256_srli_epi64(vq, 32);
9
10    // vab = [ (a_0 * b_0 = ab_0ab_1), (a_2 * b_2 = ab_2ab_3), (a_4 * b_4 =
11        ab_4ab_5), (a_6 * b_6 = ab_6ab_7) ]
12    __m256i vab = _mm256_mul_epu32(va, vb);
13
14    // vqp = [ (q_0 * p_0 = qp_0qp_1), (q_2 * p_2 = qp_2qp_3), (q_4 * p_4 =
15        qp_4qp_5), (q_6 * p_6 = qp_6qp_7) ]
16    __m256i vqp = _mm256_mul_epu32(vq, vp);
17
18    // vc = [ (ab_0 - qp_0 = c_0c_1), (ab_2 - qp_2 = c_2c_3), (ab_4 - qp_4 =
19        c_4c_5), (ab_6 - qp_6 = c_6c_7) ]
20    __m256i vc = _mm256_sub_epi64(vab, vqp);
21
22    // if (c >= p) c -= p
23    return _mm256_min_epu32(vc, _mm256_sub_epi32(vc, vp));
24 }
25
26 Vector shoup_scale_avx2(Parameters param, Vector v)
27 {
28     ulong size = v.size;
29     Vector res = init_vector(size);
30     __m256i vb = _mm256_set1_epi64x(param.b);
31     __m256i vb_precomp = _mm256_set1_epi64x(param.b_precomp);
32     __m256i vp = _mm256_set1_epi64x(param.p);
33     ulong i = 0;
34     for (; i + 7 < size; i += 8)
35     {
36         // Load [ a_0, a_1, a_2, a_3, a_4, a_5, a_6, a_7 ] = [ (a_1, a_0),
37             (a_3, a_2), (a_5, a_4), (a_7, a_6) ] but only even index will be
38             directly handled
39         __m256i even_va = _mm256_loadu_si256((__m256i const*)(v.elements +
40             i));
41
42         // Need to move odd index to even index to be handled: [ (0, a_1),
43             (0, a_3), (0, a_5), (0, a_7) ]
44         __m256i odd_va = _mm256_srli_epi64(even_va, 32);
45
46         // Process Even Lanes
47         __m256i even_vc = _shoup_avx2(even_va, vb, vb_precomp, vp);
48
49         // Process Odd Lanes:
50         __m256i odd_vc = _shoup_avx2(odd_va, vb, vb_precomp, vp);
51
52         // Shift odd by 32 bits to the left results back and merge: [ (c_1,
53             0), (c_3, 0), (c_5, 0), (c_7, 0) ]
54         __m256i odd_shifted = _mm256_slli_epi64(odd_vc, 32);
```

```

46         // Merge even and odd: [ c_0, c_1, c_2, c_3, c_4, c_5, c_6, c_7 ]
47         __m256i merge = _mm256_or_si256(even_vc, odd_shifted);
48
49         _mm256_storeu_si256((__m256i *)(res.elements + i), merge);
50     }
51     for (; i < size; i++)
52         *(res.elements + i) = shoup(*(v.elements + i), param.b,
53                                     param.b_precomp, param.p);
54     return res;
55 }

```

Listing 2: Variant 2 (AVX2)

```

1  static inline __m256i _shoup_mullo_avx2(__m256i va, __m256i vb, __m256i
2  vb_precomp, __m256i vp)
3  {
4      // vq = [ (a_0 * b_precomp_0 = q_0q_1), (a_2 * b_precomp_2 = q_2q_3),
5              (a_4 * b_precomp_4 = q_4q_5), (a_6 * b_precomp_6 = q_6q_7) ]
6      __m256i vq = _mm256_mul_epu32(va, vb_precomp);
7
8      // vq = [ (0, q_0), (0, q_2), (0, q_4), (0, q_6) ]
9      vq = _mm256_srli_epi64(vq, 32);
10
11      // Preserves only the least significant bit: vab = [ (a_0 * b_0 mod 2^32
12                = ab_0), (a_2 * b_2 mod 2^32 = ab_2), (a_4 * b_4 mod 2^32 = ab_4),
13                (a_6 * b_6 mod 2^32 = ab_6) ]
14      __m256i vab = _mm256_mullo_epi32(va, vb);
15
16      // vqp = [ [q_0 * p_0 mod 2^32 = qp_0], [q_2 * p_2 mod 2^32 = qp_2],
17                [q_4 * p_4 mod 2^32 = qp_4], [q_6 * p_6 mod 2^32 = qp_6] ]
18      __m256i vqp = _mm256_mullo_epi32(vq, vp);
19
20      // vc = [ (0, ab_0 - qp_0), (0, ab_2 - qp_2), (0, ab_4 - qp_4), (0, ab_6
21                - qp_6) ]
22      __m256i vc = _mm256_sub_epi32(vab, vqp);
23
24      // if (c >= p) c -= p
25      return _mm256_min_epu32(vc, _mm256_sub_epi32(vc, vp));
26  }
27
28  Vector shoup_scale_mullo_avx2(Parameters param, Vector v)
29  {
30      ulong size = v.size;
31      Vector res = init_vector(size);
32      __m256i vb = _mm256_set1_epi64x(param.b);
33      __m256i vb_precomp = _mm256_set1_epi64x(param.b_precomp);
34      __m256i vp = _mm256_set1_epi64x(param.p);
35      ulong i = 0;
36      for (; i + 31 < size; i += 32)
37      {
38          __m256i even_va_0 = _mm256_loadu_si256((__m256i const *) (v.elements
39              + i));
40          __m256i even_va_1 = _mm256_loadu_si256((__m256i const *) (v.elements
41              + i + 8));
42          __m256i even_va_2 = _mm256_loadu_si256((__m256i const *) (v.elements
43              + i + 16));
44          __m256i even_va_3 = _mm256_loadu_si256((__m256i const *) (v.elements
45              + i + 24));

```

```

37     __m256i odd_va_0 = _mm256_srli_epi64(even_va_0, 32);
38     __m256i odd_va_1 = _mm256_srli_epi64(even_va_1, 32);
39     __m256i odd_va_2 = _mm256_srli_epi64(even_va_2, 32);
40     __m256i odd_va_3 = _mm256_srli_epi64(even_va_3, 32);
41
42     __m256i even_vc_0 = _shoup_mullo_avx2(even_va_0, vb, vb_precomp, vp);
43     __m256i even_vc_1 = _shoup_mullo_avx2(even_va_1, vb, vb_precomp, vp);
44     __m256i even_vc_2 = _shoup_mullo_avx2(even_va_2, vb, vb_precomp, vp);
45     __m256i even_vc_3 = _shoup_mullo_avx2(even_va_3, vb, vb_precomp, vp);
46
47     __m256i odd_vc_0 = _shoup_mullo_avx2(odd_va_0, vb, vb_precomp, vp);
48     __m256i odd_vc_1 = _shoup_mullo_avx2(odd_va_1, vb, vb_precomp, vp);
49     __m256i odd_vc_2 = _shoup_mullo_avx2(odd_va_2, vb, vb_precomp, vp);
50     __m256i odd_vc_3 = _shoup_mullo_avx2(odd_va_3, vb, vb_precomp, vp);
51
52     __m256i odd_shifted_0 = _mm256_slli_epi64(odd_vc_0, 32);
53     __m256i odd_shifted_1 = _mm256_slli_epi64(odd_vc_1, 32);
54     __m256i odd_shifted_2 = _mm256_slli_epi64(odd_vc_2, 32);
55     __m256i odd_shifted_3 = _mm256_slli_epi64(odd_vc_3, 32);
56
57     __m256i merge_0 = _mm256_or_si256(even_vc_0, odd_shifted_0);
58     __m256i merge_1 = _mm256_or_si256(even_vc_1, odd_shifted_1);
59     __m256i merge_2 = _mm256_or_si256(even_vc_2, odd_shifted_2);
60     __m256i merge_3 = _mm256_or_si256(even_vc_3, odd_shifted_3);
61
62     _mm256_storeu_si256((__m256i*)(res.elements + i), merge_0);
63     _mm256_storeu_si256((__m256i*)(res.elements + i + 8), merge_1);
64     _mm256_storeu_si256((__m256i*)(res.elements + i + 16), merge_2);
65     _mm256_storeu_si256((__m256i*)(res.elements + i + 24), merge_3);
66 }
67 for (; i < size; i++)
68     *(res.elements + i) = shoup(*(v.elements + i), param.b,
69     param.b_precomp, param.p);
69 return res;
70 }

```

Listing 3: Variant 3 (AVX2)

```

1 static inline __m256i _mulhi_emul_avx2(__m256i va, __m256i vb)
2 {
3     // res_even = [ (a_0 * b_0), (a_2 * b_2), (a_4 * b_4), (a_6 * b_6) ]
4     __m256i res_even = _mm256_mul_epu32(va, vb);
5
6     // res_odd = [ (a_1 * b_0), (a_3 * b_2), (a_5 * b_4), (a_7 * b_6) ]
7     __m256i res_odd = _mm256_mul_epu32(_mm256_srli_epi64(va, 32), vb);
8
9     // hi_even = [ (0, hi(a_0 * b_0)), (0, hi(a_2 * b_2)), (0, hi(a_4 *
10     // b_4)), (0, hi(a_6 * b_6)) ]
11     __m256i hi_even = _mm256_srli_epi64(res_even, 32);
12
13     // Removes the least significant bits: hi_odd = [ (hi(a_1 * b_0), 0),
14     // (hi(a_3 * b_2), 0), (hi(a_5 * b_4), 0), (hi(a_7 * b_6), 0) ]
15     __m256i hi_odd = _mm256_and_si256(res_odd,
16     _mm256_set1_epi64x(0xFFFFFFFF00000000ULL));
17
18     // [ (hi(a_1 * b_0), hi(a_0 * b_0)), (hi(a_3 * b_2), hi(a_2 * b_2)),
19     // (hi(a_5 * b_4), hi(a_4 * b_4)), (hi(a_7 * b_6), hi(a_6 * b_6)) ]
20     return _mm256_or_si256(hi_even, hi_odd);
21 }

```

```

19 static inline __m256i _shoup_mullo_v2_avx2(__m256i va, __m256i vb, __m256i
    vb_precomp, __m256i vp)
20 {
21     __m256i vq = _mulhi_emul_avx2(va, vb_precomp);
22     __m256i vab = _mm256_mullo_epi32(va, vb);
23     __m256i vqp = _mm256_mullo_epi32(vq, vp);
24     __m256i vc = _mm256_sub_epi32(vab, vqp);
25     return _mm256_min_epu32(vc, _mm256_sub_epi32(vc, vp));
26 }
27
28 Vector shoup_scale_mullo_v2_avx2(Parameters param, Vector v)
29 {
30     ulong size = v.size;
31     Vector res = init_vector(size);
32     __m256i vb = _mm256_set1_epi32(param.b);
33     __m256i vb_precomp = _mm256_set1_epi32(param.b_precomp);
34     __m256i vp = _mm256_set1_epi32(param.p);
35     ulong i = 0;
36     for (; i + 31 < size; i += 32)
37     {
38         __m256i va_0 = _mm256_loadu_si256((__m256i const *)(v.elements + i));
39         __m256i va_1 = _mm256_loadu_si256((__m256i const *)(v.elements + i +
40             8));
41         __m256i va_2 = _mm256_loadu_si256((__m256i const *)(v.elements + i +
42             16));
43         __m256i va_3 = _mm256_loadu_si256((__m256i const *)(v.elements + i +
44             24));
45         __m256i vc_0 = _shoup_mullo_v2_avx2(va_0, vb, vb_precomp, vp);
46         __m256i vc_1 = _shoup_mullo_v2_avx2(va_1, vb, vb_precomp, vp);
47         __m256i vc_2 = _shoup_mullo_v2_avx2(va_2, vb, vb_precomp, vp);
48         __m256i vc_3 = _shoup_mullo_v2_avx2(va_3, vb, vb_precomp, vp);
49         _mm256_storeu_si256((__m256i *)(res.elements + i), vc_0);
50         _mm256_storeu_si256((__m256i *)(res.elements + i + 8), vc_1);
51         _mm256_storeu_si256((__m256i *)(res.elements + i + 16), vc_2);
52         _mm256_storeu_si256((__m256i *)(res.elements + i + 24), vc_3);
53     }
54     for (; i < size; i++)
55         *(res.elements + i) = shoup(*(v.elements + i), param.b,
56             param.b_precomp, param.p);
57     return res;
58 }

```

A.2 AVX-512

Listing 4: Variant 1 (AVX-512)

```

1 static inline __m512i _shoup_avx512(__m512i va, __m512i vb, __m512i
    vb_precomp, __m512i vp)
2 {
3     __m512i vq = _mm512_mul_epu32(va, vb_precomp);
4     vq = _mm512_srli_epi64(vq, 32);
5     __m512i vab = _mm512_mul_epu32(va, vb);
6     __m512i vqp = _mm512_mul_epu32(vq, vp);
7     __m512i vc = _mm512_sub_epi64(vab, vqp);
8     return _mm512_min_epu32(vc, _mm512_sub_epi32(vc, vp));
9 }
10
11 Vector shoup_scale_avx512(Parameters param, Vector v)

```

```

12 {
13     ulong size = v.size;
14     Vector res = init_vector(size);
15     __m512i vb = _mm512_set1_epi64(param.b);
16     __m512i vb_precomp = _mm512_set1_epi64(param.b_precomp);
17     __m512i vp = _mm512_set1_epi64(param.p);
18     ulong i = 0;
19     for (; i + 15 < size; i += 16)
20     {
21         __m512i even_va = _mm512_loadu_si512((__m512i const *)(v.elements +
22             i));
23
24         __m512i odd_va = _mm512_srli_epi64(even_va, 32);
25
26         __m512i even_vc = _shoup_avx512(even_va, vb, vb_precomp, vp);
27
28         __m512i odd_vc = _shoup_avx512(odd_va, vb, vb_precomp, vp);
29
30         __m512i odd_shifted = _mm512_slli_epi64(odd_vc, 32);
31
32         __m512i merge = _mm512_or_si512(even_vc, odd_shifted);
33
34         _mm512_storeu_si512((__m512 *) (res.elements + i), merge);
35     }
36     for (; i < size; i++)
37         *(res.elements + i) = shoup(*(v.elements + i), param.b,
38             param.b_precomp, param.p);
39     return res;
40 }

```

Listing 5: Variant 2 (AVX-512)

```

1 static inline __m512i _shoup_mullo_avx512(__m512i va, __m512i vb, __m512i
2     vb_precomp, __m512i vp)
3 {
4     __m512i vq = _mm512_mul_epu32(va, vb_precomp);
5     vq = _mm512_srli_epi64(vq, 32);
6     __m512i vab = _mm512_mullo_epi32(va, vb);
7     __m512i vqp = _mm512_mullo_epi32(vq, vp);
8     __m512i vc = _mm512_sub_epi32(vab, vqp);
9     return _mm512_min_epu32(vc, _mm512_sub_epi32(vc, vp));
10 }
11 Vector shoup_scale_mullo_avx512(Parameters param, Vector v)
12 {
13     ulong size = v.size;
14     Vector res = init_vector(size);
15     __m512i vb = _mm512_set1_epi64(param.b);
16     __m512i vb_precomp = _mm512_set1_epi64(param.b_precomp);
17     __m512i vp = _mm512_set1_epi64(param.p);
18     ulong i = 0;
19     for (; i + 63 < size; i += 64)
20     {
21         __m512i even_va_0 = _mm512_loadu_si512((__m512i const *)(v.elements
22             + i));
23         __m512i even_va_1 = _mm512_loadu_si512((__m512i const *)(v.elements
24             + i + 16));
25         __m512i even_va_2 = _mm512_loadu_si512((__m512i const *)(v.elements
26             + i + 32));
27         __m512i even_va_3 = _mm512_loadu_si512((__m512i const *)(v.elements

```

```

25         + i + 48));
26     __m512i odd_va_0 = _mm512_srli_epi64(even_va_0, 32);
27     __m512i odd_va_1 = _mm512_srli_epi64(even_va_1, 32);
28     __m512i odd_va_2 = _mm512_srli_epi64(even_va_2, 32);
29     __m512i odd_va_3 = _mm512_srli_epi64(even_va_3, 32);
30
31     __m512i even_vc_0 = _shoup_mullo_avx512(even_va_0, vb, vb_precomp,
32     vp);
33     __m512i even_vc_1 = _shoup_mullo_avx512(even_va_1, vb, vb_precomp,
34     vp);
35     __m512i even_vc_2 = _shoup_mullo_avx512(even_va_2, vb, vb_precomp,
36     vp);
37     __m512i even_vc_3 = _shoup_mullo_avx512(even_va_3, vb, vb_precomp,
38     vp);
39
40     __m512i odd_vc_0 = _shoup_mullo_avx512(odd_va_0, vb, vb_precomp, vp);
41     __m512i odd_vc_1 = _shoup_mullo_avx512(odd_va_1, vb, vb_precomp, vp);
42     __m512i odd_vc_2 = _shoup_mullo_avx512(odd_va_2, vb, vb_precomp, vp);
43     __m512i odd_vc_3 = _shoup_mullo_avx512(odd_va_3, vb, vb_precomp, vp);
44
45     __m512i odd_shifted_0 = _mm512_slli_epi64(odd_vc_0, 32);
46     __m512i odd_shifted_1 = _mm512_slli_epi64(odd_vc_1, 32);
47     __m512i odd_shifted_2 = _mm512_slli_epi64(odd_vc_2, 32);
48     __m512i odd_shifted_3 = _mm512_slli_epi64(odd_vc_3, 32);
49
50     __m512i merge_0 = _mm512_or_si512(even_vc_0, odd_shifted_0);
51     __m512i merge_1 = _mm512_or_si512(even_vc_1, odd_shifted_1);
52     __m512i merge_2 = _mm512_or_si512(even_vc_2, odd_shifted_2);
53     __m512i merge_3 = _mm512_or_si512(even_vc_3, odd_shifted_3);
54
55     _mm512_storeu_si512((__m512 *) (res.elements + i), merge_0);
56     _mm512_storeu_si512((__m512 *) (res.elements + i + 16), merge_1);
57     _mm512_storeu_si512((__m512 *) (res.elements + i + 32), merge_2);
58     _mm512_storeu_si512((__m512 *) (res.elements + i + 48), merge_3);
59 }
60
61 for (; i < size; i++)
62     *(res.elements + i) = shoup(*(v.elements + i), param.b,
63     param.b_precomp, param.p);
64
65 return res;
66 }

```

Listing 6: Variant 3 (AVX-512)

```

1 static inline __m512i _mulhi_emul_avx512(__m512i va, __m512i vb)
2 {
3     __m512i res_even = _mm512_mul_epu32(va, vb);
4     __m512i res_odd = _mm512_mul_epu32(_mm512_srli_epi64(va, 32), vb);
5     __m512i hi_even = _mm512_srli_epi64(res_even, 32);
6     __m512i hi_odd = _mm512_and_si512(res_odd,
7     _mm512_set1_epi64(0xFFFFFFFF00000000ULL));
8     return _mm512_or_si512(hi_even, hi_odd);
9 }
10
11 static inline __m512i _shoup_mullo_v2_avx512(__m512i va, __m512i vb, __m512i
12 vb_precomp, __m512i vp)
13 {
14     __m512i vq = _mulhi_emul_avx512(va, vb_precomp);
15     __m512i vab = _mm512_mullo_epi32(va, vb);
16     __m512i vqp = _mm512_mullo_epi32(vq, vp);

```



```

15     __m512i vc = _mm512_sub_epi32(vab, vqp);
16     return _mm512_min_epu32(vc, _mm512_sub_epi32(vc, vp));
17 }
18
19 Vector shoup_scale_mullo_v2_avx512(Parameters param, Vector v)
20 {
21     ulong size = v.size;
22     Vector res = init_vector(size);
23     __m512i vb = _mm512_set1_epi32(param.b);
24     __m512i vb_precomp = _mm512_set1_epi32(param.b_precomp);
25     __m512i vp = _mm512_set1_epi32(param.p);
26     ulong i = 0;
27     for (; i + 63 < size; i += 64)
28     {
29         __m512i va_0 = _mm512_loadu_si512((__m512i const *)(v.elements + i));
30         __m512i va_1 = _mm512_loadu_si512((__m512i const *)(v.elements + i +
31         16));
32         __m512i va_2 = _mm512_loadu_si512((__m512i const *)(v.elements + i +
33         32));
34         __m512i va_3 = _mm512_loadu_si512((__m512i const *)(v.elements + i +
35         48));
36
37         __m512i vc_0 = _shoup_mullo_v2_avx512(va_0, vb, vb_precomp, vp);
38         __m512i vc_1 = _shoup_mullo_v2_avx512(va_1, vb, vb_precomp, vp);
39         __m512i vc_2 = _shoup_mullo_v2_avx512(va_2, vb, vb_precomp, vp);
40         __m512i vc_3 = _shoup_mullo_v2_avx512(va_3, vb, vb_precomp, vp);
41
42         _mm512_storeu_si512((__m512i *)(res.elements + i), vc_0);
43         _mm512_storeu_si512((__m512i *)(res.elements + i + 16), vc_1);
44         _mm512_storeu_si512((__m512i *)(res.elements + i + 32), vc_2);
45         _mm512_storeu_si512((__m512i *)(res.elements + i + 48), vc_3);
46     }
47     for (; i < size; i++)
48         *(res.elements + i) = shoup(*(v.elements + i), param.b,
49         param.b_precomp, param.p);
50     return res;
51 }

```

A.3 NEON

Listing 7: Variant 1 (NEON)

```
1 static inline uint32x2_t _shoup_neon(uint32x2_t va, uint32x2_t vb,  
   uint32x2_t vb_precomp, uint32x2_t vp)  
2 {  
3     // vab_precomp = [ a_0 * b_precomp_0 = ab_precomp_0, a_1 * b_precomp_1 =  
   ab_precomp_1 ]  
4     uint64x2_t vab_precomp = vmull_u32(va, vb_precomp);  
5  
6     // vq = [ q_0, q_1 ]  
7     uint32x2_t vq = vshrn_n_u64(vab_precomp, 32);  
8  
9     // vab = [ a_0 * b_0 = ab_0, a_1 * b_1 = ab_1 ]  
10    uint64x2_t vab = vmull_u32(va, vb);  
11  
12    // vqp = [ q_0 * p_0 = qp_0, q_1 * p_1 = qp_1 ]  
13    uint64x2_t vqp = vmull_u32(vq, vp);  
14  
15    // vc = [ ab_0 - qp_0 = c_0, ab_1 - qp_1 = c_1 ]  
16    uint32x2_t vc = vmovn_u64(vsubq_u64(vab, vqp));  
17  
18    // if (c >= p) c -= p  
19    return vmin_u32(vc, vsub_u32(vc, vp));  
20 }  
21  
22 Vector shoup_scale_neon(Parameters param, Vector v)  
23 {  
24     ulong size = v.size;  
25     Vector res = init_vector(size);  
26     ulong n = size - (size % 4);  
27     uint32x2_t vb = vdup_n_u32(param.b);  
28     uint32x2_t vp = vdup_n_u32(param.p);  
29     uint32x2_t vb_precomp = vdup_n_u32(param.b_precomp);  
30     ulong i = 0;  
31     for (; i + 1 < n; i += 2)  
32     {  
33         // Load [ a_0, a_1 ]  
34         uint32x2_t va = vld1_u32(v.elements + i);  
35  
36         // Stocks the value  
37         vst1_u32(res.elements + i, _shoup_neon(va, vb, vb_precomp, vp));  
38     }  
39     for (; i < size; i++)  
40         *(res.elements + i) = shoup(*(v.elements + i), param.b,  
   param.b_precomp, param.p);  
41     return res;  
42 }
```

Listing 8: Variant 2 (NEON)

```
1 static inline uint32x2_t _shoup_mullo_neon(uint32x2_t va, uint32x2_t vb,  
   uint32x2_t vb_precomp, uint32x2_t vp)  
2 {  
3     // vab_precomp = [ a_0 * b_precomp_0 = ab_precomp_0, a_1 * b_precomp_1 =  
   ab_precomp_1 ]  
4     uint64x2_t vab_precomp = vmull_u32(va, vb_precomp);  
5  
6     // vq = [ q_0, q_1 ]  
7     uint32x2_t vq = vshrn_n_u64(vab_precomp, 32);
```

```

8
9 // vab = [ a_0 * b_0 mod 2^32 = ab_0, a_1 * b_1 mod 2^32 = ab_1 ]
10 uint32x2_t vab = vmul_u32(va, vb);
11
12 // vqp = [ q_0 * p_0 mod 2^32 = qp_0, q_1 * p_1 mod 2^32 = qp_1 ]
13 uint32x2_t vqp = vmul_u32(vq, vp);
14
15 // vc = [ ab_0 - qp_0 = c_0, ab_1 - qp_1 = c_1 ]
16 uint32x2_t vc = vsub_u32(vab, vqp);
17 return vmin_u32(vc, vsub_u32(vc, vp));
18 }
19
20 Vector shoup_scale_mullo_neon(Parameters param, Vector v)
21 {
22     ulong size = v.size;
23     Vector res = init_vector(size);
24     ulong n = size - (size % 4);
25     uint32x2_t vb = vdup_n_u32(param.b);
26     uint32x2_t vp = vdup_n_u32(param.p);
27     uint32x2_t vb_precomp = vdup_n_u32(param.b_precomp);
28     ulong i = 0;
29     for (; i + 7 < n; i += 8)
30     {
31         uint32x2_t va_0 = vld1_u32(v.elements + i);
32         uint32x2_t va_1 = vld1_u32(v.elements + i + 2);
33         uint32x2_t va_2 = vld1_u32(v.elements + i + 4);
34         uint32x2_t va_3 = vld1_u32(v.elements + i + 6);
35
36         vst1_u32(res.elements + i, _shoup_mullo_neon(va_0, vb, vb_precomp,
37             vp));
38         vst1_u32(res.elements + i + 2, _shoup_mullo_neon(va_1, vb,
39             vb_precomp, vp));
40         vst1_u32(res.elements + i + 4, _shoup_mullo_neon(va_2, vb,
41             vb_precomp, vp));
42         vst1_u32(res.elements + i + 6, _shoup_mullo_neon(va_3, vb,
43             vb_precomp, vp));
44     }
45     for (; i < size; i++)
46         *(res.elements + i) = shoup(*(v.elements + i), param.b,
47             param.b_precomp, param.p);
48     return res;
49 }

```